

ASSUMPTION

Docs are similar if they have ~~the~~
~~SAME~~ words

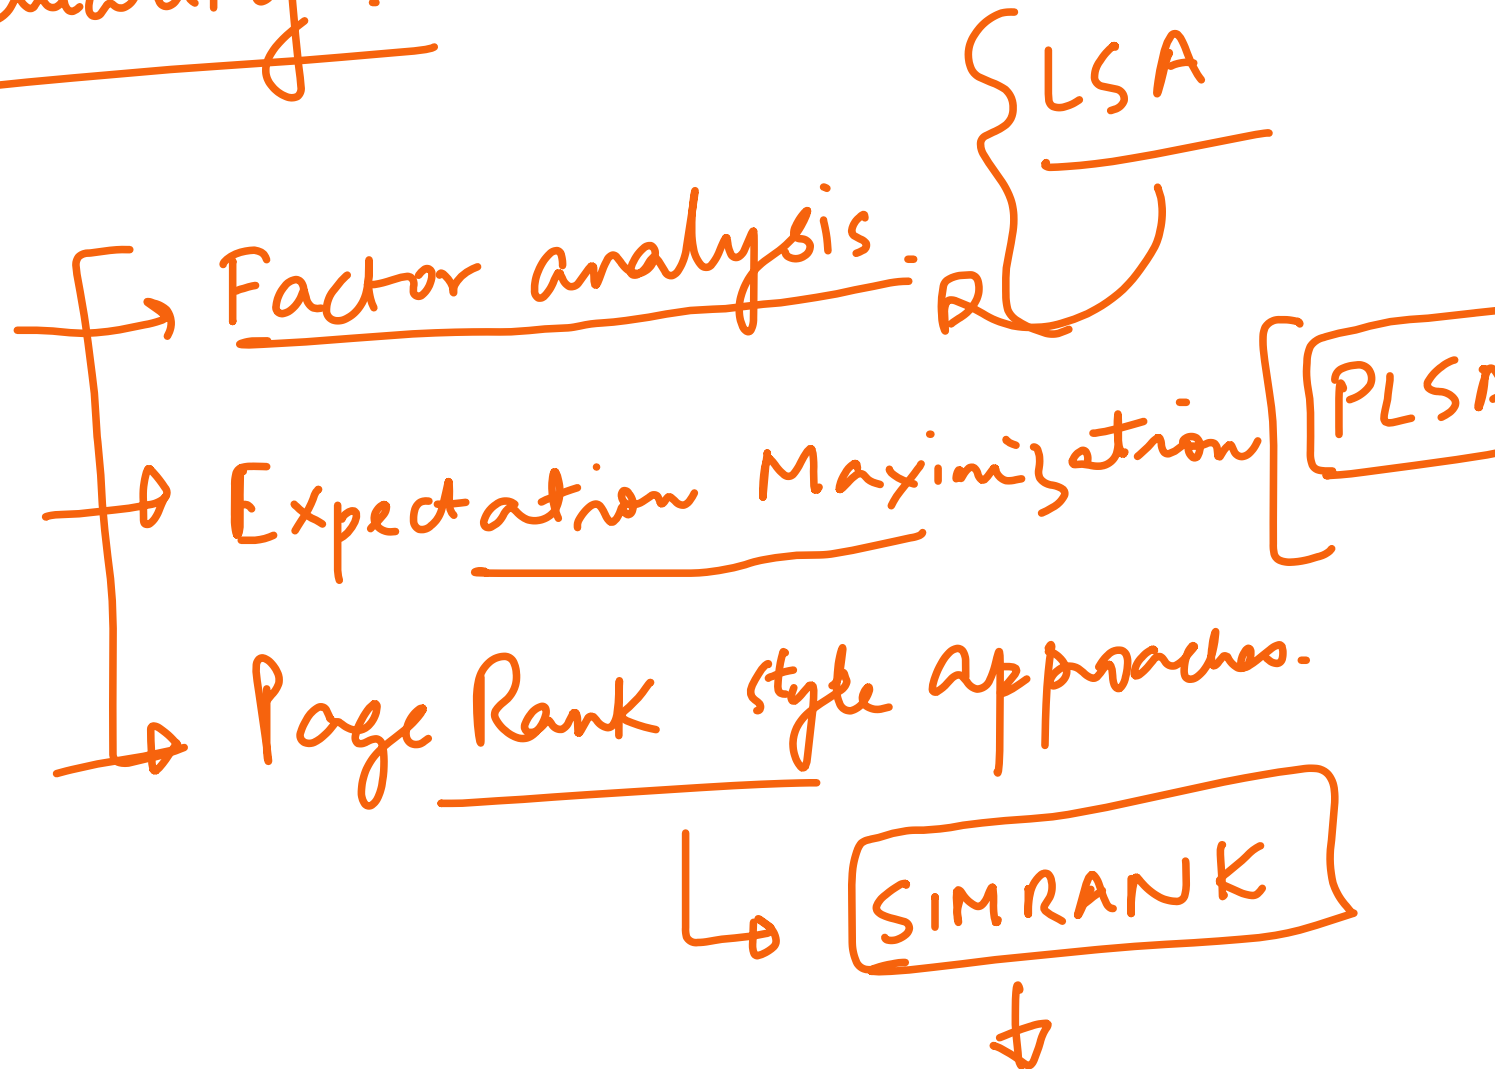
SIMILAR

Words are similar if they occur in ~~the~~
~~SAME~~ docs.

SIMILAR

CIRCULARITY

Circularity.




PageRank

$$P(d|s) \propto \underbrace{P(q|d)}_{\text{PageRank.}} \boxed{P(d)}$$


↑ PageRank.

A page is important
if it is pointed to by several important
pages.



Latent Semantic Analysis.

Factor Analysis.



Stretching
rotating.

$$\textcircled{M} \vec{q} = \vec{q'}$$

operator

COVARIANCE MATRIX

Words: cat, dog, chair

$$\textcircled{M} \rightarrow \begin{matrix} \text{cat} & \text{dog} & \text{chair} \\ \begin{bmatrix} \text{cat} & \text{dog} & \text{chair} \\ 1 & 0.8 & 0 \\ 0.8 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix} \end{matrix} \left. \vphantom{\begin{matrix} \text{cat} & \text{dog} & \text{chair} \\ 1 & 0.8 & 0 \\ 0.8 & 1 & 0 \\ 0 & 0 & 1 \end{matrix}} \right\} \text{Principal Component Analysis}$$

query

$$\begin{bmatrix} 1 & 0 & 0 \end{bmatrix}$$

cat dog chair

$$M \vec{d} = \vec{d'}$$

each concept
is a linear combination
of terms

UNIMODE FACTOR ANALYSIS.

rotation & stretching



vector sum of stretches.

$$M \vec{q}$$



$$\vec{v}_1, \vec{v}_2, \vec{v}_3$$

eigenvectors

$$\left. \begin{aligned} M \vec{v}_1 &= \lambda_1 \vec{v}_1 \\ M \vec{v}_2 &= \lambda_2 \vec{v}_2 \\ M \vec{v}_3 &= \lambda_3 \vec{v}_3 \end{aligned} \right\}$$

dimensionality
reduction.

$$\vec{q} = \alpha_1 \vec{v}_1 + \alpha_2 \vec{v}_2 + \alpha_3 \vec{v}_3$$

$$\begin{aligned} \vec{q} = M \vec{q} &= \alpha_1 M \vec{v}_1 + \alpha_2 M \vec{v}_2 + \alpha_3 M \vec{v}_3 \quad \times \\ &= \alpha_1 \lambda_1 \vec{v}_1 + \alpha_2 \lambda_2 \vec{v}_2 + \alpha_3 \lambda_3 \vec{v}_3 \dots \end{aligned}$$

1.00	0.8	0.00003	0
0	0	0	0

Latent dimensions.

taste

	sour	semi-sweet	sweet
Green	0.8		0.1
yellow		0.9	
red	0.1		0.8

A hand-drawn diagram illustrating a data structure. On the left, a vertical arrow points downwards, labeled "data". To the right of the arrow is a large square bracket structure. Above the top of this bracket is a wavy line labeled "word".

unripe ripe

$$\begin{matrix}
 & \begin{matrix} \text{ur} & \text{sr} & \text{r} \end{matrix} \\
 \begin{matrix} \text{g} \\ \text{g} \\ \text{r} \end{matrix} & \left[\begin{array}{ccc} & & \\ & & \\ & & \end{array} \right]
 \end{matrix}$$

colour \times ripeness

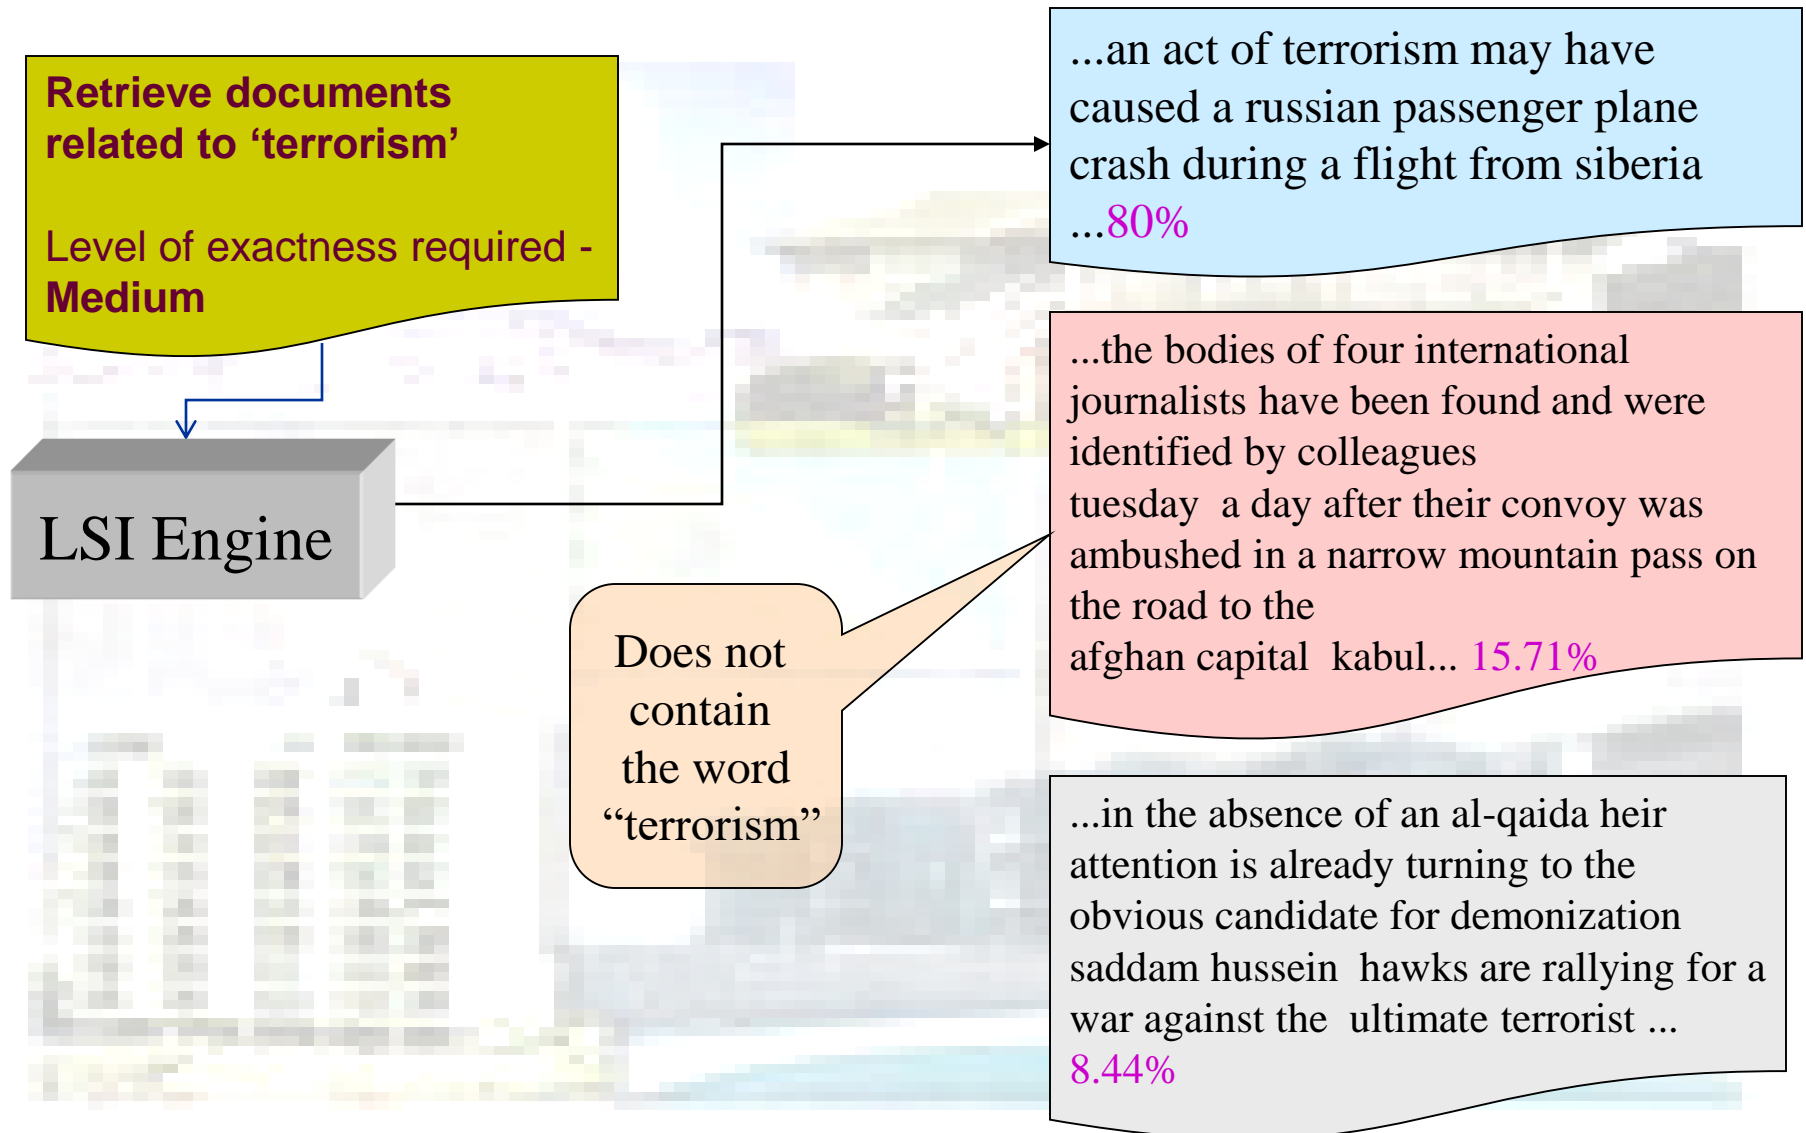
colour x ripeness ripeness x taste

concept

$\text{doc} \begin{bmatrix} \quad \end{bmatrix} \text{con} \begin{bmatrix} \quad \end{bmatrix}$
 $\text{graph} \begin{bmatrix} \quad \end{bmatrix}$
 $\underbrace{m \times k + k \times n}_{k(m+n)}$

Latent Semantic Indexing

Search using LSI - an example



Latent Semantic Analysis

Starting point : a term document matrix

Document #	filtering	indexing	clustering	decompose	sediments	purification	oscillation	matrix	differential
1	1	1	1	0	0	0	0	1	0
2	1	0	1	0	0	0	0	0	0
3	1	1	1	0	0	0	0	0	0
4	0	0	0	1	1	1	0	0	0
5	0	0	0	1	0	1	0	0	0
6	1	0	0	1	1	1	0	0	0
7	0	0	0	0	0	0	1	1	1
8	0	0	0	0	0	0	1	0	1
9	0	0	0	1	0	0	1	1	1

Objective : Determine the best set of underlying factors or concepts, that best explain the relationship between terms and documents

Assumption : There is some underlying or latent structure in word usage that is partially obscured by variability in word choice

Procedure : Singular Value Decomposition

Key Idea : Factor Analysis

A bit of Linear Algebra

- In Linear Algebra, a matrix like the term-document matrix can be viewed as an **Operator**
- The matrix can act upon a vector (when it is multiplied with that vector), and relocate it to a different position.

$$M = \begin{pmatrix} 2 & -1 & 1 \\ -1 & 2 & -1 \\ 1 & -1 & 2 \end{pmatrix}$$

can act on the vector

$$\vec{A} = \begin{pmatrix} 1 \\ 2 \\ 3 \end{pmatrix}$$

and move it to a new location given by $M \vec{A}$:

$$M \vec{A} = \begin{pmatrix} 3 \\ 0 \\ 5 \end{pmatrix}$$

The action of a matrix can be viewed as a combination of translation and rotation of vector

Goal: characterizing a matrix formally in terms of its properties that govern its action on vectors

The concept of eigenvectors does precisely that

Eigenvectors

We consider all vectors \vec{x} that, when acted on by M , stretch themselves to a different location $\lambda \vec{x}$, where λ is a scalar, but do not undergo any rotation. Thus

$$M \vec{x} = \lambda \vec{x} \tag{1}$$

The vectors satisfying (1) are called eigenvectors, and each of these eigenvectors is associated with a corresponding value of λ referred to as an eigenvalue.

$(M - \lambda I) \vec{x} = 0$, where I is an identity matrix of dimensions matching M ; this is called the characteristic equation.

in our example, we have the following three eigenvectors

$$v_1 = \begin{pmatrix} -1 \\ 0 \\ 1 \end{pmatrix}, v_2 = \begin{pmatrix} 1 \\ 1 \\ 0 \end{pmatrix} \text{ and } v_3 = \begin{pmatrix} 1 \\ -1 \\ 1 \end{pmatrix}$$

associated with the eigenvalues $\lambda_1 = 1$, $\lambda_2 = 1$ and $\lambda_3 = 4$ respectively.

Why are eigenvectors important?

We now study the effect of M on any arbitrary vector \vec{x}

$$\vec{x} = \begin{pmatrix} 4 \\ -1 \\ 4 \end{pmatrix}$$

We can express \vec{x} as a linear combination of v_1, v_2 and v_3 .

The revised position $M \vec{x}$ is now given by

$$\begin{aligned} M \vec{x} &= M (1v_1 + 2v_2 + 3v_3) \\ &= Mv_1 + 2Mv_2 + 3Mv_3 \\ &= \lambda_1 v_1 + 2\lambda_2 v_2 + 3\lambda_3 v_3 \end{aligned}$$

The interesting aspect of this rewrite is that we can see that the total effect of M on \vec{x} is expressed as a weighted combination of effects due to each eigenvector.

Eigenvectors having very small eigenvalues associated with them have a small effect on the operation of M on \vec{x} .

Critical Intuition behind SVD/LSI

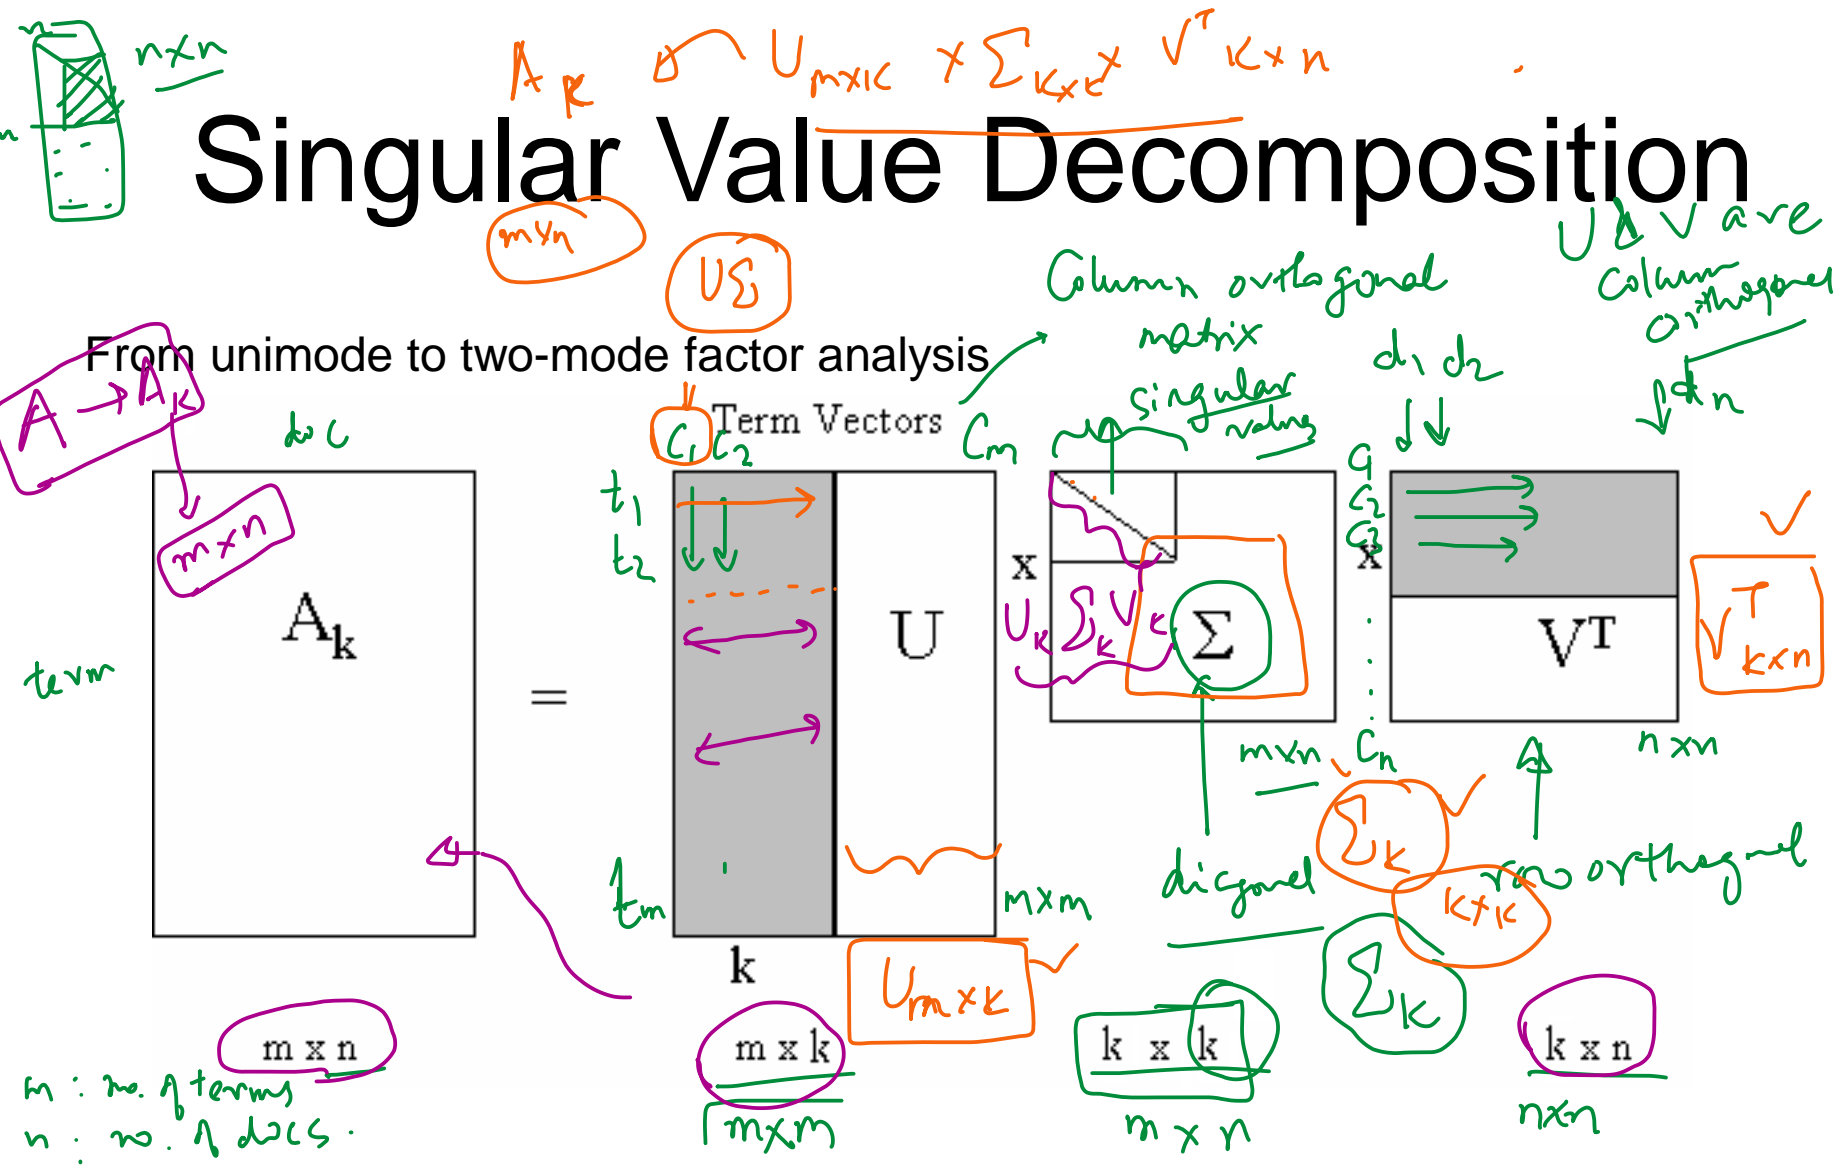
Matrix Diagonalization Theorem

For a given square real valued $m \times m$ matrix M with linearly independent eigenvectors, we can obtain a factorization

$$\underline{M = U \Lambda U^{-1}} \quad \left[\begin{array}{ccc} \cdot & & \\ & \ddots & \\ & & \cdot \end{array} \right]$$

such that the columns of U are the eigenvectors of M , and Λ is a diagonal matrix whose diagonal elements are eigenvalues of M arranged in decreasing order. This result is due to the Matrix Diagonalization Theorem. This result applies to square matrices, but not to rectangular ones like the term-document matrix.


Singular Value Decomposition



A_k is the reduced rank approximation of the original matrix A

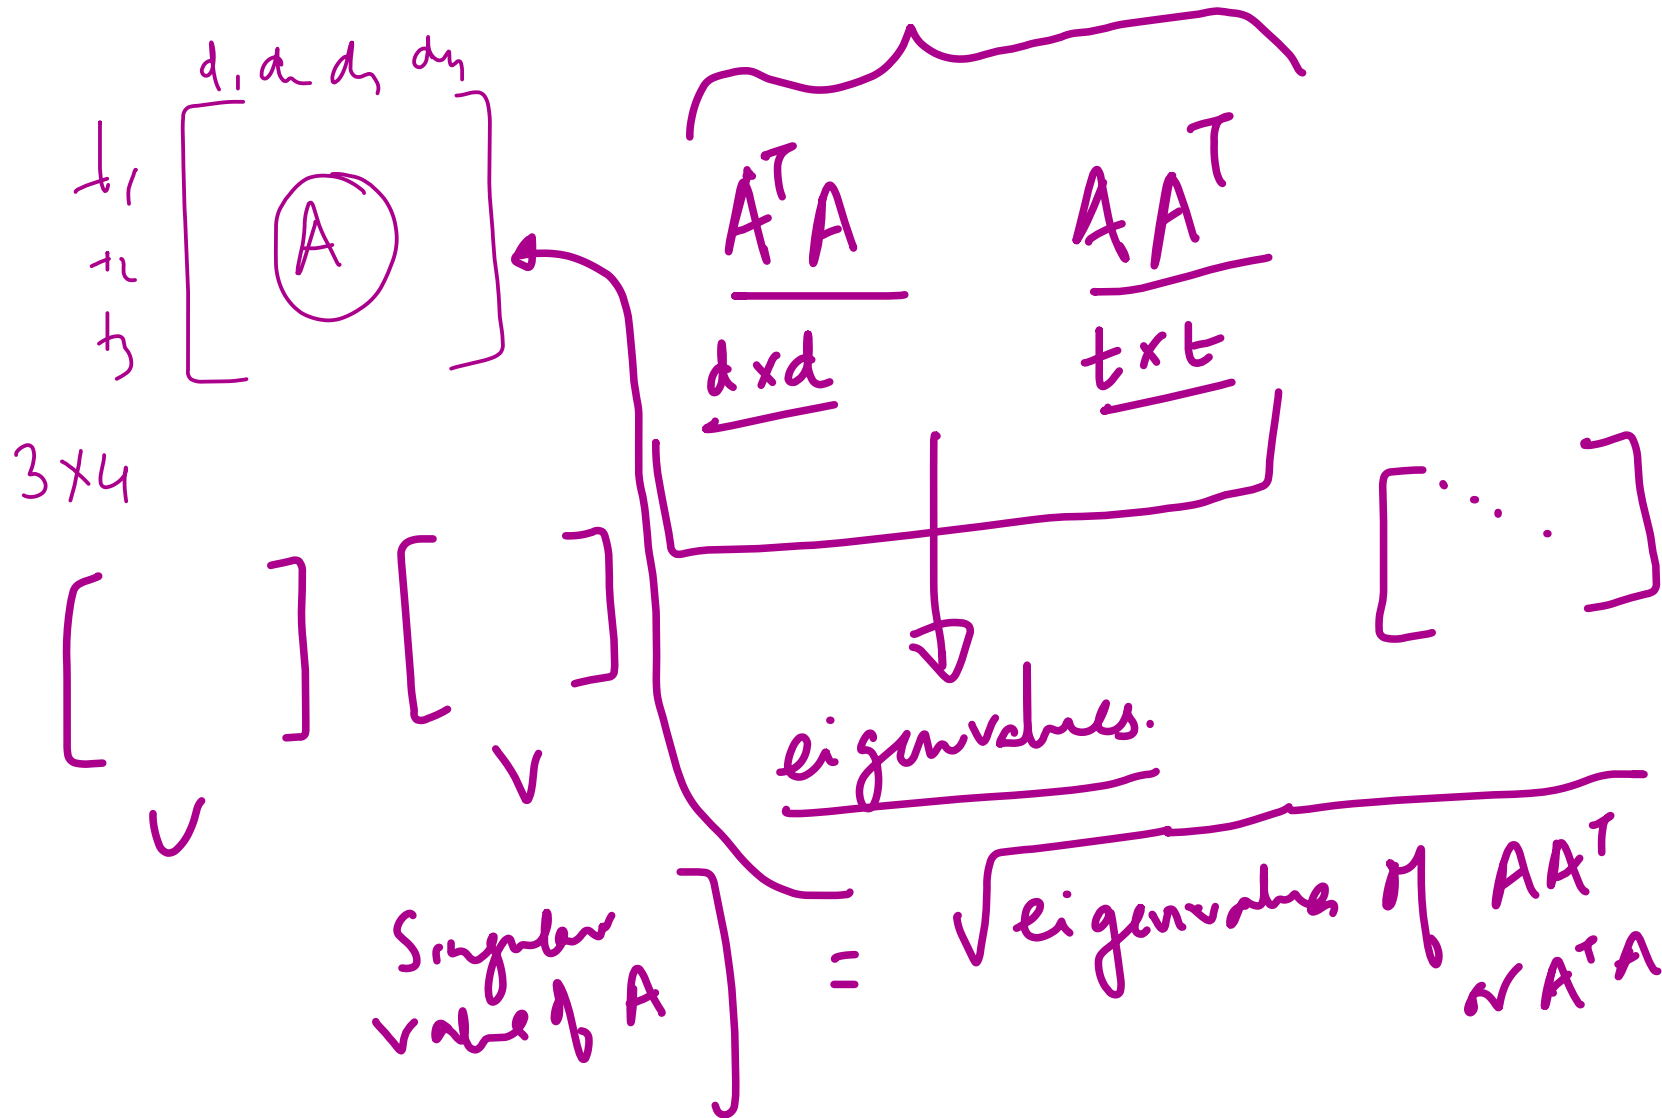
Singular Value Decomposition

From unimode to two-mode factor analysis



A_k is the reduced rank approximation of the original matrix A

Connection of SVD to Matrix Diagonalization



Example


$k=3$

IR

Chemistry

Maths

Document #	filtering	indexing	clustering	decompose	sediments	purification	oscillation	matrix	differential
1	1	1	1	0	0	0	0	1	0
2	1	0	1	0	0	0	0	0	0
3	1	1	1	0	0	0	0	0	0
4	0	0	0	1	1	1	0	0	0
5	0	0	0	1	0	1	0	0	0
6	1	0	0	1	1	1	0	0	0
7	0	0	0	0	0	0	1	1	1
8	0	0	0	0	0	0	1	0	1
9	0	0	0	1	0	0	1	1	1



Document #	filtering	indexing	clustering	decompose	sediments	purification	oscillation	matrix	differential
1	1.17	0.83	1.11	-0.00	-0.00	-0.06	0.16	0.59	0.17
2	0.89	0.53	0.72	0.00	0.06	0.04	-0.1	0.22	-0.1
3	1.07	0.73	0.98	-0.05	0.04	-0.01	-0.11	0.32	-0.11
4	0.24	-0.12	-0.14	1.09	0.74	1.01	0.00	-0.04	0.00
5	0.16	-0.1	-0.11	0.82	0.54	0.74	0.04	-0.00	0.04
6	0.67	0.15	0.23	1.16	0.82	1.09	-0.06	0.06	-0.06
7	-0.02	0.07	0.05	0.23	-0.12	-0.15	1.02	0.84	1.017
8	-0.13	-0.03	-0.07	0.18	-0.08	-0.09	0.74	0.56	0.74
9	0.04	0.02	-0.02	0.67	0.15	0.23	1.11	0.89	1.11

Example 2

Technical Memo Example

Titles


- | | |
|-----|--|
| c1: | <i>Human machine interface for Lab ABC computer applications</i> |
| c2: | <i>A survey of user opinion of computer system response time</i> |
| c3: | <i>The EPS user interface management system</i> |
| c4: | <i>System and human system engineering testing of EPS</i> |
| c5: | <i>Relation of user-perceived response time to error measurement</i> |
| m1: | <i>The generation of random, binary, unordered trees</i> |
| m2: | <i>The intersection graph of paths in trees</i> |
| m3: | <i>Graph minors IV: Widths of trees and well-quasi-ordering</i> |
| m4: | <i>Graph minors: A survey</i> |

Terms

Documents

[illegible]

Example 2



des .

A =

1	0	0	1	0	0	0	0	0
1	0	1	0	0	0	0	0	0
1	1	0	0	0	0	0	0	0
0	1	1	0	1	0	0	0	0
0	1	1	2	0	0	0	0	0
0	1	0	0	1	0	0	0	0
0	1	0	0	1	0	0	0	0
0	0	1	1	0	0	0	0	0
0	1	0	0	0	0	0	0	1
0	0	0	0	0	1	1	1	0
0	0	0	0	0	0	1	1	1
0	0	0	0	0	0	0	1	1

U =

0.22	-0.11	0.29	-0.41	-0.11	-0.34	0.52	-0.06	-0.41
0.20	-0.07	0.14	-0.55	0.28	0.50	-0.07	-0.01	-0.11
0.24	0.04	-0.16	-0.59	-0.11	-0.25	-0.30	0.06	0.49
0.40	0.06	-0.34	0.10	0.33	0.38	0.00	0.00	0.01
0.64	-0.17	0.36	0.33	-0.16	-0.21	-0.17	0.03	0.27
0.27	0.11	-0.43	0.07	0.08	-0.17	0.28	-0.02	-0.05
0.27	0.11	-0.43	0.07	0.08	-0.17	0.28	-0.02	-0.05
0.30	-0.14	0.33	0.19	0.11	0.27	0.03	-0.02	-0.17
0.21	0.27	-0.18	-0.03	-0.54	0.08	-0.47	-0.04	-0.58
0.01	0.49	0.23	0.03	0.59	-0.39	-0.29	0.25	-0.23
0.04	0.62	0.22	0.00	-0.07	0.11	0.16	-0.68	0.23
0.03	0.45	0.14	-0.01	-0.30	0.28	0.34	0.68	0.18

Σ =

3.34	2.54	2.35	1.64	1.50	1.31	0.85	0.56	0.36
------	------	------	------	------	------	------	------	------

V =

0.20	-0.06	0.11	-0.95	0.05	-0.08	0.18	-0.01	-0.06
0.61	0.17	-0.50	-0.03	-0.21	-0.26	-0.43	0.05	0.24
0.46	-0.03	0.21	0.04	0.38	0.72	-0.24	0.01	0.02
0.54	-0.23	0.57	0.27	-0.21	-0.37	0.26	-0.02	-0.08
0.28	0.11	-0.51	0.15	0.33	0.03	0.67	-0.06	-0.26
0.00	0.19	0.10	0.02	0.39	-0.30	-0.34	0.45	-0.62
0.01	0.44	0.19	0.02	0.35	-0.21	-0.15	-0.76	0.02
0.02	0.62	0.25	0.01	0.15	0.00	0.25	0.45	0.51
0.08	0.53	0.08	-0.03	-0.60	0.36	-0.04	-0.07	-0.45

Example 2

U_2

Σ_2

V_2'

0.22	-0.11
0.20	-0.07
0.24	0.04
0.40	0.06
0.64	-0.17
0.27	0.11
0.27	0.11
0.30	-0.14
0.21	0.27
0.01	0.49
0.04	0.62
0.03	0.45

3.34
2.54

0.20	0.61	0.46	0.54	0.28	0.00	0.02	0.02	0.08
-0.06	0.17	-0.13	-0.23	0.11	0.19	0.44	0.62	0.53

$$A_2 = U_2 \Sigma_2 V_2' =$$

0.16	0.40	0.38	0.47	0.18	-0.05	-0.12	-0.16	-0.09
0.14	0.37	0.33	0.40	0.16	-0.03	-0.07	-0.10	-0.04
0.15	0.51	0.36	0.41	0.24	0.02	0.06	0.09	0.12
0.26	0.84	0.61	0.70	0.39	0.03	0.08	0.12	0.19
0.45	1.23	1.05	1.27	0.56	-0.07	-0.15	-0.21	-0.05
0.16	0.58	0.38	0.42	0.28	0.06	0.13	0.19	0.22
0.16	0.58	0.38	0.42	0.28	0.06	0.13	0.19	0.22
0.22	0.55	0.51	0.63	0.24	-0.07	-0.14	-0.20	-0.11
0.10	0.53	0.23	0.21	0.27	0.14	0.31	0.44	0.42
-0.06	0.23	-0.14	-0.27	0.14	0.24	0.55	0.77	0.66
-0.06	0.34	-0.15	-0.30	0.20	0.31	0.69	0.98	0.85
-0.04	0.25	-0.10	-0.21	0.15	0.22	0.50	0.71	0.62

$A =$

1	0	0	1	0	0	0	0	0
1	0	1	0	0	0	0	0	0
1	1	0	0	0	0	0	0	0
0	1	1	0	1	0	0	0	0
0	1	1	2	0	0	0	0	0
0	1	0	0	1	0	0	0	0
0	1	0	0	1	0	0	0	0
0	0	1	1	0	0	0	0	0
0	1	0	0	0	0	0	0	1
0	0	0	0	0	1	1	1	0
0	0	0	0	0	0	1	1	1
0	0	0	0	0	0	0	1	1

The “dual mode” advantage



Things we can do

- Compute document similarities
- Compute term similarities
- Compute relevance of a term to a document

$$\begin{array}{l} A_k \\ t \times d \end{array} \quad \begin{array}{l} / \\ A_k A_k' \end{array}$$


←

$$\begin{array}{l} A_k' A_k \end{array}$$

←

Term-term similarities

	filtering	indexing	clustering	decompose	sediments	purification	oscillation	matrix	differential
filtering	4	2	3	1	1	1	0	1	0
indexing	2	2	2	0	0	0	0	1	0
clustering	3	2	3	0	0	0	0	1	0
decompose	1	0	0	4	2	3	1	1	1
sediments	1	0	0	2	2	2	0	0	0
purification	1	0	0	3	2	3	0	0	0
oscillation	0	0	0	1	0	0	3	2	3
matrix	1	1	1	1	0	0	2	3	2
differential	0	0	0	1	0	0	3	2	3



	filtering	indexing	clustering	decompose	sediments	purification	oscillation	matrix	differential
filtering	3.86	2.28	3.10	1.12	0.93	1.07	-0.13	1.21	-0.11
indexing	2.28	1.56	2.09	-0.05	0.04	-0.07	0.06	0.91	0.07
clustering	3.10	2.09	2.80	-0.04	0.10	-0.03	-0.04	1.13	-0.03
decompose	1.12	-0.05	-0.04	3.74	2.26	3.08	1.08	0.90	1.08
sediments	0.93	0.04	0.10	2.26	1.56	2.10	-0.05	0.03	-0.05
purification	1.07	-0.07	-0.03	3.08	2.10	2.84	-0.01	0.02	-0.01
oscillation	-0.13	0.06	-0.04	1.08	-0.05	-0.01	2.87	2.29	2.87
matrix	1.21	0.91	1.13	0.90	0.03	0.02	2.29	2.32	2.30
differential	-0.11	0.07	-0.03	1.08	-0.05	-0.01	2.87	2.30	2.87


Refer this link for Geometry of SVD

<http://websites.uwlax.edu/twill/svd/perpframes/index.html>

$A \vec{x}$

Hanger (Column Orthogonal)

Before the Hit




\vec{p}_1 \vec{p}_2

$$\begin{pmatrix} 0.36 & -0.93 \\ 0.93 & 0.36 \end{pmatrix}$$

U

Aligner

Before the Hit




$$\begin{pmatrix} 0.36 & 0.93 \\ -0.93 & 0.36 \end{pmatrix} \vec{x}$$

$$V^T$$

Stretcher


Before the Hit



$$\begin{pmatrix} 3 & 0 \\ 0 & 2 \end{pmatrix}$$


$$SVD = H_{\text{anger}} S_{\text{tretcher}} A_{\text{ligner}}$$

Key point: The action of a matrix product: $A=BC$ is the same as the action of the C followed by the action of B .



$$\begin{bmatrix} & \\ & \end{bmatrix} \begin{bmatrix} & \\ & \end{bmatrix} \rightarrow \begin{bmatrix} & \\ & \end{bmatrix}$$

$f \times t$



TERM
CONCEPT

$$d \begin{bmatrix} & \\ & \end{bmatrix} \begin{bmatrix} q \\ & \end{bmatrix} =$$

$t \times d$ $d \times 1$

$d \times t$ $t \times 1$

$$H \begin{bmatrix} S & A & q \end{bmatrix}$$

\downarrow
 V^T


LSA

- Idea (Deerwester et al):


“We would like a representation in which a set of terms, which by itself is incomplete and unreliable evidence of the relevance of a given document, is replaced by some other set of entities which are more reliable indicants. We take advantage of the implicit higher-order (or latent) structure in the association of terms and documents to reveal such relationships.”

ML \leftrightarrow Info Theory



SVD for compression



$$y = a + b \cdot x$$



✓ $M \subset Kary$



To send the matrix M you need to send $256 \times 264 = 67584$ numbers.

To send the rank 36 approximation to M you need only send
the first 36 singular values,

the first 36 hanger vectors, each of which has 256 entries,

the first 36 aligner vectors, each of which has 264 entries.


So in total you need to send only $36(1+256+264)=18756$ numbers

<http://websites.uwlax.edu/twill/svd/compression/index.html>

The Right Choice of Dimensions?

- Reduce dimensionality to:
 - Fewer dimensions, more “collapsing of axes”, better recall, worse precision
 - More dimensions, less collapsing, worse recall, better precision

- Cross validation over training data?




Mapping Queries to New Space

$$\hat{q} = \underbrace{q' U_k \Sigma_k^{-1}}_I$$

[Is LSA better for Synonymy
or polysemy?]


LSA and Orders of Co-Occurrence

Doc 1




(a)

Doc 1




Doc 2




(b)


Doc 1




Doc 2




Doc 3



(c)



(d)



(e)

Empirical

SVD : an important property

It can be shown that \hat{M} is the best k-rank approximation to M in the least-squares sense. The quality of an approximation M_A is measured by the Frobenius Norm of the “discrepancy” matrix $X = M - M_A$, which is given by:



$$\checkmark \quad \|X\|_F = \sqrt{\sum_{i=1}^m \sum_{j=1}^n X_{ij}^2}.$$


$$\begin{bmatrix} \cdot \\ \cdot \\ \cdot \end{bmatrix}_A \quad \begin{bmatrix} \cdot \\ \cdot \\ \cdot \end{bmatrix}_{A_K}$$

SVD and rank reduction

Intuition : Distort a matrix to generate an approximation that has lower rank than the original

Refer: <http://www.ams.org/publicoutreach/feature-column/fcarc-svd>





$$\sigma_1 = 14.15$$

$$\sigma_2 = 4.67$$

$$\sigma_3 = 3.00$$

$$\sigma_4 = 0.21$$

$$\sigma_5 = 0.19$$

...

$$\sigma_{15} = 0.05$$

$$M \approx \mathbf{u}_1 \sigma_1 \mathbf{v}_1^T + \mathbf{u}_2 \sigma_2 \mathbf{v}_2^T + \mathbf{u}_3 \sigma_3 \mathbf{v}_3^T$$

$$L = 0.9 C_1 + 0.05 C_2 + 0.05 C_3$$

$$\sigma_1 = 14.15$$

$$\sigma_2 = 4.67$$

$$\sigma_3 = 3.00$$


$$\sigma_4 = 0.21$$

$$\sigma_5 = 0.19$$


...

$$\sigma_{15} = 0.05$$


$$M \approx \mathbf{u}_1 \sigma_1 \mathbf{v}_1^T + \mathbf{u}_2 \sigma_2 \mathbf{v}_2^T + \mathbf{u}_3 \sigma_3 \mathbf{v}_3^T$$



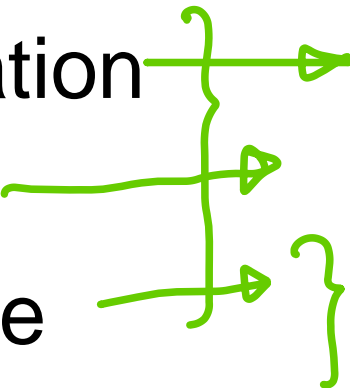
Noisy image




Improved image



Update Techniques

- Recomputation
 - Folding In
 - SVD Update
- 
- Hand-drawn green annotations on the list items:
- A vertical curly brace on the right side of the first three items, spanning from 'Recomputation' to 'SVD Update'.
 - A horizontal arrow pointing right from the end of 'Recomputation'.
 - A horizontal arrow pointing right from the end of 'Folding In'.
 - A horizontal arrow pointing right from the end of 'SVD Update'.
 - A vertical curly brace on the right side of the 'SVD Update' item.

Is LSI good for classification?



LSA Applications

- Information Retrieval : Term Weighting
- Cross language Retrieval
- Other matrices: Matching people instead of documents

How would you compute SVD of
a 3×4 matrix by hand?

References

- Indexing by Latent Semantic Analysis by Scott Deerwester, Susan T. Dumais, George W. Furnas, Thomas K. Landauer, Richard Harshman